

PCOS modeling

Erick Navarro

2023-01-25

Contents

Creating the classifiers	1
Setup	1
Load packages and data	1
Data splitting	1
Imputation on training set	2
Define the models	10
Logistic regression modeling	10
Elastic net regression	12
Random Forest	19
Model comparison	19
Effect of class imbalance on the best model (if time allows)	19

Creating the classifiers

This report contains the details of the generation and comparison of several classifiers whose aim is to classify people with and without Polycystic ovary syndrome (PCOS). The goal of this project is to develop and validate a model that takes easily collected variables, and test its performance against models that possess variables obtained with increasing invasive procedures. For a summary of the results and a clearer picture of the project, please read the Final_report.pdf file in this repository.

Setup

Load packages and data

Data splitting

After loading the data and packages, we will proceed to split the data into training and validation sets.

```
data = data %>%
  column_to_rownames("id")

train.index <- caret::createDataPartition(data$pcos, p = .7, list=FALSE)

train <- data[ train.index,]
valid <- data[-train.index,]

#Check the negative/positive ration in the validation set
table(train$pcos)
```

```
##
## No Yes
## 255 124
```

```
table(train$pcos)[1]/ table(train$pcos)[2]
```

```
## No
## 2.056452
```

```
#Check the negative/positive ration in the validation set
table(valid$pcos)
```

```
##
## No Yes
## 109 53
```

```
table(valid$pcos)[1]/ table(valid$pcos)[2]
```

```
## No
## 2.056604
```

```
#Check the negative/positive ration in the original data set
table(data$pcos)[1]/ table(data$pcos)[2]
```

```
## No
## 2.056497
```

Imputation on training set

Following the data splitting, we will impute missing values. As we observed in the EDA analysis, the proportion of missing data is overall very low, which makes all the variables with missingness suitable for imputation

```
# Explore missingness in training set
sapply(train, function(x) sum(is.na(x)))
```

```
##          pcos          age          weight          height          bmi
##          0          0          0          0          0
##    blood_group    pulse_rate          rr          hb          cycle
##          0          2          0          0          0
##    cycle_length marriage_status    pregnant no_of_abortions    i_betahcg
##          0          0          0          0          0
##    ii_betahcg          fsh          lh    fsh_lh_ratio          hip
##          0          0          1          1          0
##          waist waist_hip_ratio          tsh          amh          prl
##          0          0          0          2          0
##          vitd3          prg          rbs    weight_gain    hair_growth
##          1          0          0          0          0
##    skin_darkening    hair_loss    pimples    fast_food    reg_exercise
##          0          0          0          1          0
##    bp_systolic    bp_diastolic    follicle_no_l    follicle_no_r    avg_f_size_l
##          0          0          0          0          0
##    avg_f_size_r    endometrium
##          0          0
```

For imputation, we will use the mice R package, a widely used software to handle missing data. We will use the default options, which use the most appropriate methodology depending on the class of the data to be imputed, which are the following: “By default, the method uses pmm, predictive mean matching (numeric data) logreg, logistic regression imputation (binary data, factor with 2 levels) polyreg, polytomous regression imputation for un- ordered categorical data (factor > 2 levels) polr, proportional odds model for (ordered, > 2 levels).”


```
chained_train = mice::mice(train)
```

```
##
##  iter imp variable
##    1  1 pulse_rate lh fsh_lh_ratio amh vitd3 fast_food
##    1  2 pulse_rate lh fsh_lh_ratio amh vitd3 fast_food
##    1  3 pulse_rate lh fsh_lh_ratio amh vitd3 fast_food
##    1  4 pulse_rate lh fsh_lh_ratio amh vitd3 fast_food
##    1  5 pulse_rate lh fsh_lh_ratio amh vitd3 fast_food
##    2  1 pulse_rate lh fsh_lh_ratio amh vitd3 fast_food
##    2  2 pulse_rate lh fsh_lh_ratio amh vitd3 fast_food
##    2  3 pulse_rate lh fsh_lh_ratio amh vitd3 fast_food
##    2  4 pulse_rate lh fsh_lh_ratio amh vitd3 fast_food
##    2  5 pulse_rate lh fsh_lh_ratio amh vitd3 fast_food
##    3  1 pulse_rate lh fsh_lh_ratio amh vitd3 fast_food
##    3  2 pulse_rate lh fsh_lh_ratio amh vitd3 fast_food
##    3  3 pulse_rate lh fsh_lh_ratio amh vitd3 fast_food
##    3  4 pulse_rate lh fsh_lh_ratio amh vitd3 fast_food
##    3  5 pulse_rate lh fsh_lh_ratio amh vitd3 fast_food
##    4  1 pulse_rate lh fsh_lh_ratio amh vitd3 fast_food
##    4  2 pulse_rate lh fsh_lh_ratio amh vitd3 fast_food
##    4  3 pulse_rate lh fsh_lh_ratio amh vitd3 fast_food
##    4  4 pulse_rate lh fsh_lh_ratio amh vitd3 fast_food
##    4  5 pulse_rate lh fsh_lh_ratio amh vitd3 fast_food
##    5  1 pulse_rate lh fsh_lh_ratio amh vitd3 fast_food
##    5  2 pulse_rate lh fsh_lh_ratio amh vitd3 fast_food
##    5  3 pulse_rate lh fsh_lh_ratio amh vitd3 fast_food
```


```
## 5 4 pulse_rate lh fsh_lh_ratio amh vitd3 fast_food
## 5 5 pulse_rate lh fsh_lh_ratio amh vitd3 fast_food
```

```
## Warning: Number of logged events: 98
```


```
#Explore the imputed data and check that the generated value are plausible
stripplot(chained_train, pulse_rate, pch = 19, xlab = "Imputation number")
```



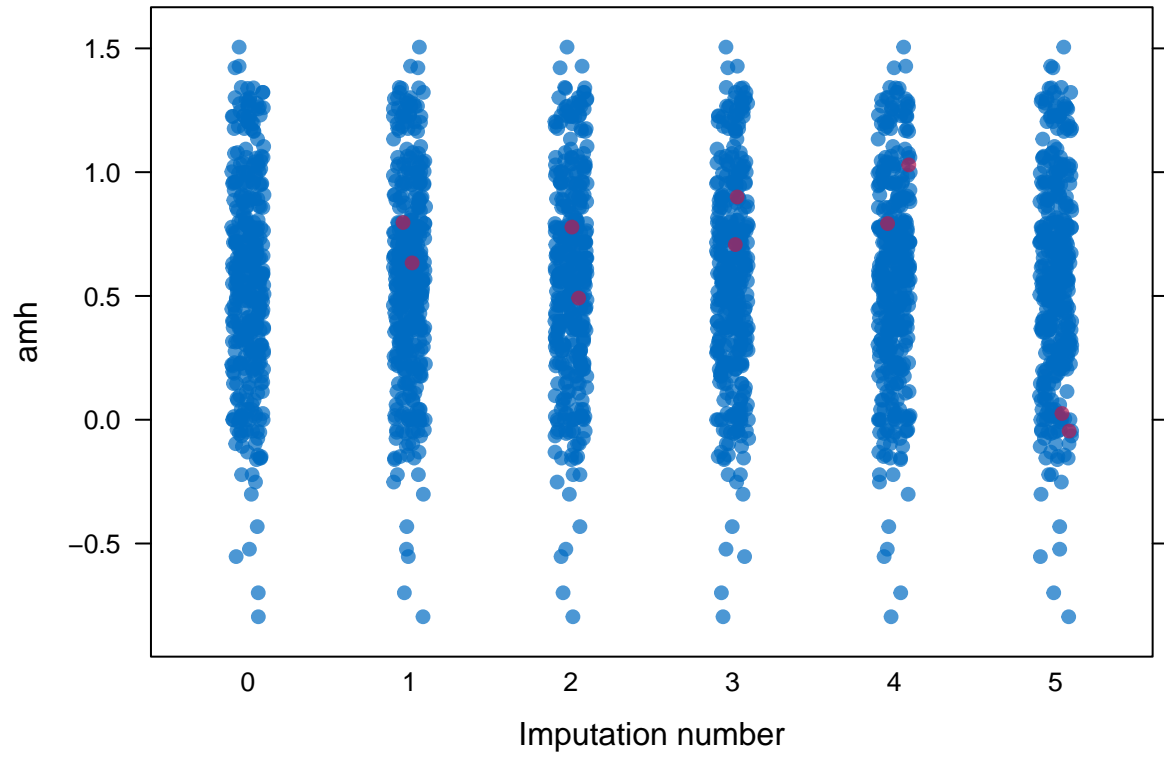
```
stripplot(chained_train, lh, pch = 19, xlab = "Imputation number")
```




```
stripplot(chained_train, fsh_lh_ratio, pch = 19, xlab = "Imputation number")
```




```
stripplot(chained_train, amh, pch = 19, xlab = "Imputation number")
```



```
stripplot(chained_train, vitd3, pch = 19, xlab = "Imputation number")
```



```
stripplot(chained_train, fast_food, pch = 19, xlab = "Imputation number")
```

#It looks fine so we will proceed to extract the imputed data

```
chained_train = complete(data=chained_train)

sapply(chained_train, function(x) sum(is.na(x)))
```

```
##          pcos          age          weight          height          bmi
##           0           0           0           0           0
##   blood_group   pulse_rate          rr          hb          cycle
##           0           0           0           0           0
##   cycle_length marriage_status   pregnant no_of_abortions   i_betahcg
##           0           0           0           0           0
##     ii_betahcg          fsh          lh   fsh_lh_ratio          hip
##           0           0           0           0           0
##       waist waist_hip_ratio          tsh          amh          prl
##           0           0           0           0           0
##       vitd3          prg          rbs   weight_gain   hair_growth
##           0           0           0           0           0
## skin_darkening   hair_loss   pimples   fast_food   reg_exercise
##           0           0           0           0           0
##   bp_systolic bp_diastolic follicle_no_l follicle_no_r   avg_f_size_l
##           0           0           0           0           0
##   avg_f_size_r   endometrium
##           0           0
```

Define the models

For the aim of this project, we will develop 4 different models with an increasing number of variables according to its easiness to collect. Then, model 1 will contain only variables that are collected through the patient history, model 2 will have clinical examination ones added, model 3 blood tests and model 4 all the relevant variables available in the dataset.

It is worth mentioning that we removed some non-relevant variables from the dataset based on the scientific literature of the field. A more detailed information about this selection can be found in the final report available in this repository.

```
# Create dataset 1 using only variables obtained through patient history
modell1_vars = c("pcos","age","cycle","cycle_length",
                "no_of_abortions", "weight_gain", "hair_growth", "skin_darkening",
                "hair_loss","pimples","fast_food", "reg_exercise")

# Create dataset 2 using variables obtained through patient history and clinical examination
modell2_vars = c(modell1_vars, "weight","height","bmi",
                "hip", "waist","waist_hip_ratio",
                "bp_systolic", "bp_diastolic")

# Create dataset 3 using variables obtained through patient history, clinical examination and blood tests
modell3_vars = c(modell2_vars, "fsh", "lh", "fsh_lh_ratio",
                "amh", "prl", "vitd3", "prg", "rbs")

# Create dataset 4 using variables obtained through patient history, clinical examination, blood tests
modell4_vars = c(modell3_vars, "follicle_no_l", "follicle_no_r" ,
                "avg_f_size_l", "avg_f_size_r", "endometrium")
```

Logistic regression modeling

```
fitControl <- trainControl(
  method = 'cv',
  number = 5,
  savePredictions = 'final',
  classProbs = TRUE,
  summaryFunction=twoClassSummary)

set.seed(504)
#Since we have few samples, I will use bootstrapping instead of cross fold validation

models_logreg = list()

for (model in list(modell1_vars, modell2_vars, modell3_vars, modell4_vars)){
  cv_model = caret::train(
    pcos ~ .,
    data = chained_train %>%
      select(all_of(model)),
    method = "glm",
    family = "binomial",
    trControl = fitControl)
  models_logreg = append(models_logreg, list(cv_model))
}
```

```
## Warning in train.default(x, y, weights = w, ...): The metric "Accuracy" was not
## in the result set. ROC will be used instead.
```

```
## Warning in train.default(x, y, weights = w, ...): The metric "Accuracy" was not
## in the result set. ROC will be used instead.
```

```
## Warning in train.default(x, y, weights = w, ...): The metric "Accuracy" was not
## in the result set. ROC will be used instead.
```

```
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

```
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

```
## Warning in train.default(x, y, weights = w, ...): The metric "Accuracy" was not
## in the result set. ROC will be used instead.
```

```
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

```
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

```
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

```
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

```
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

```
names(models_logreg) = c("model1", "model2", "model3", "model4")
```

```
models_logreg
```

```
## $model1
## Generalized Linear Model
##
## 379 samples
## 11 predictor
## 2 classes: 'No', 'Yes'
##
## No pre-processing
## Resampling: Cross-Validated (5 fold)
## Summary of sample sizes: 304, 303, 303, 303, 303
## Resampling results:
##
## ROC      Sens      Spec
## 0.9050523 0.9019608 0.718
##
## $model2
## Generalized Linear Model
##
## 379 samples
## 19 predictor
## 2 classes: 'No', 'Yes'
##
```

```
## No pre-processing
## Resampling: Cross-Validated (5 fold)
## Summary of sample sizes: 303, 304, 303, 303, 303
## Resampling results:
##
##      ROC          Sens          Spec
##  0.8943791  0.9019608  0.6946667
##
##
## $model3
## Generalized Linear Model
##
## 379 samples
## 27 predictor
## 2 classes: 'No', 'Yes'
##
## No pre-processing
## Resampling: Cross-Validated (5 fold)
## Summary of sample sizes: 303, 304, 303, 303, 303
## Resampling results:
##
##      ROC          Sens          Spec
##  0.8718627  0.8941176  0.6943333
##
##
## $model4
## Generalized Linear Model
##
## 379 samples
## 32 predictor
## 2 classes: 'No', 'Yes'
##
## No pre-processing
## Resampling: Cross-Validated (5 fold)
## Summary of sample sizes: 303, 304, 303, 303, 303
## Resampling results:
##
##      ROC          Sens          Spec
##  0.9350327  0.9137255  0.815
```

Elastic net regression

```
models_EN = list()

for (model in list(model1_vars, model2_vars, model3_vars, model4_vars)){
  en_model = caret::train(
    pcos ~ .,
    trControl = fitControl,
    data = chained_train %>%
      select(all_of(model)),
    method = "glmnet",
    tuneGrid = expand.grid(alpha = seq(0.1,.2,by = 0.05),
```

```

                                lambda = seq(0.05,0.3,by = 0.05)),
  verbose = FALSE,
  metric="ROC")
models_EN = append(models_EN, list(en_model))
}

names(models_EN) = c("model1","model2", "model3", "model4")
models_EN

```

```

## $model1
## glmnet
##
## 379 samples
## 11 predictor
## 2 classes: 'No', 'Yes'
##
## No pre-processing
## Resampling: Cross-Validated (5 fold)
## Summary of sample sizes: 303, 304, 303, 303, 303
## Resampling results across tuning parameters:
##
##  alpha  lambda  ROC      Sens      Spec
##  0.10   0.05   0.9064837 0.9254902 0.7253333
##  0.10   0.10   0.9051275 0.9372549 0.6933333
##  0.10   0.15   0.9040359 0.9490196 0.6613333
##  0.10   0.20   0.9027745 0.9529412 0.6373333
##  0.10   0.25   0.9026307 0.9568627 0.6053333
##  0.10   0.30   0.9026307 0.9568627 0.5650000
##  0.15   0.05   0.9056176 0.9254902 0.7253333
##  0.15   0.10   0.9037157 0.9411765 0.6933333
##  0.15   0.15   0.9019771 0.9490196 0.6533333
##  0.15   0.20   0.9018333 0.9568627 0.6133333
##  0.15   0.25   0.9008072 0.9568627 0.5730000
##  0.15   0.30   0.8992386 0.9568627 0.5086667
##  0.20   0.05   0.9037353 0.9215686 0.7253333
##  0.20   0.10   0.9027680 0.9411765 0.6773333
##  0.20   0.15   0.9011993 0.9529412 0.6533333
##  0.20   0.20   0.8990752 0.9568627 0.5810000
##  0.20   0.25   0.8971928 0.9568627 0.5246667
##  0.20   0.30   0.8968856 0.9568627 0.4446667
##
## ROC was used to select the optimal model using the largest value.
## The final values used for the model were alpha = 0.1 and lambda = 0.05.
##
## $model2
## glmnet
##
## 379 samples
## 19 predictor
## 2 classes: 'No', 'Yes'
##
## No pre-processing
## Resampling: Cross-Validated (5 fold)

```

```

## Summary of sample sizes: 304, 303, 303, 303, 303
## Resampling results across tuning parameters:
##
##   alpha  lambda  ROC          Sens          Spec
##   0.10   0.05    0.9011961  0.9215686  0.6930000
##   0.10   0.10    0.9031307  0.9372549  0.6770000
##   0.10   0.15    0.9031438  0.9490196  0.6610000
##   0.10   0.20    0.9033137  0.9490196  0.6366667
##   0.10   0.25    0.9041176  0.9529412  0.5876667
##   0.10   0.30    0.9044379  0.9529412  0.5636667
##   0.15   0.05    0.9012288  0.9254902  0.6933333
##   0.15   0.10    0.9044052  0.9411765  0.6853333
##   0.15   0.15    0.9045882  0.9450980  0.6526667
##   0.15   0.20    0.9050850  0.9529412  0.5880000
##   0.15   0.25    0.9044575  0.9529412  0.5636667
##   0.15   0.30    0.9034118  0.9607843  0.5316667
##   0.20   0.05    0.9015556  0.9294118  0.7013333
##   0.20   0.10    0.9039477  0.9372549  0.6770000
##   0.20   0.15    0.9052614  0.9411765  0.6363333
##   0.20   0.20    0.9042222  0.9529412  0.5956667
##   0.20   0.25    0.9026601  0.9568627  0.5556667
##   0.20   0.30    0.9014706  0.9647059  0.4590000
##
## ROC was used to select the optimal model using the largest value.
## The final values used for the model were alpha = 0.2 and lambda = 0.15.
##
## $model3
## glmnet
##
## 379 samples
## 27 predictor
## 2 classes: 'No', 'Yes'
##
## No pre-processing
## Resampling: Cross-Validated (5 fold)
## Summary of sample sizes: 304, 303, 303, 303, 303
## Resampling results across tuning parameters:
##
##   alpha  lambda  ROC          Sens          Spec
##   0.10   0.05    0.8999869  0.9058824  0.6930000
##   0.10   0.10    0.9050327  0.9176471  0.6930000
##   0.10   0.15    0.9045882  0.9294118  0.6686667
##   0.10   0.20    0.9049150  0.9372549  0.6366667
##   0.10   0.25    0.9041111  0.9450980  0.6206667
##   0.10   0.30    0.9040980  0.9529412  0.5806667
##   0.15   0.05    0.9015686  0.9058824  0.6930000
##   0.15   0.10    0.9049085  0.9176471  0.6930000
##   0.15   0.15    0.9038105  0.9294118  0.6526667
##   0.15   0.20    0.9033203  0.9411765  0.6286667
##   0.15   0.25    0.9032941  0.9529412  0.5806667
##   0.15   0.30    0.9017190  0.9529412  0.5323333
##   0.20   0.05    0.9031699  0.9098039  0.6930000
##   0.20   0.10    0.9046144  0.9215686  0.6930000
##   0.20   0.15    0.9025294  0.9333333  0.6366667


```

```


## 0.20 0.20 0.9025359 0.9529412 0.6126667
## 0.20 0.25 0.9009542 0.9529412 0.5483333
## 0.20 0.30 0.9011111 0.9568627 0.4603333
##
## ROC was used to select the optimal model using the largest value.
## The final values used for the model were alpha = 0.1 and lambda = 0.1.
##
## $model4
## glmnet
##
## 379 samples
## 32 predictor
## 2 classes: 'No', 'Yes'
##
## No pre-processing
## Resampling: Cross-Validated (5 fold)
## Summary of sample sizes: 304, 303, 303, 303, 303
## Resampling results across tuning parameters:
##
## alpha lambda ROC Sens Spec
## 0.10 0.05 0.9597255 0.9529412 0.7803333
## 0.10 0.10 0.9595621 0.9607843 0.7403333
## 0.10 0.15 0.9592288 0.9647059 0.7163333
## 0.10 0.20 0.9587451 0.9647059 0.7083333
## 0.10 0.25 0.9590784 0.9725490 0.7000000
## 0.10 0.30 0.9590784 0.9803922 0.6920000
## 0.15 0.05 0.9597320 0.9490196 0.7723333
## 0.15 0.10 0.9597190 0.9607843 0.7323333
## 0.15 0.15 0.9597124 0.9647059 0.7163333
## 0.15 0.20 0.9582745 0.9686275 0.7000000
## 0.15 0.25 0.9573203 0.9843137 0.6920000
## 0.15 0.30 0.9573268 0.9843137 0.6513333
## 0.20 0.05 0.9587647 0.9490196 0.7723333
## 0.20 0.10 0.9595556 0.9568627 0.7403333
## 0.20 0.15 0.9585948 0.9647059 0.7163333
## 0.20 0.20 0.9579673 0.9764706 0.7000000
## 0.20 0.25 0.9582876 0.9843137 0.6430000
## 0.20 0.30 0.9594052 0.9882353 0.6270000
##
## ROC was used to select the optimal model using the largest value.
## The final values used for the model were alpha = 0.15 and lambda = 0.05.

#Plot parameter tuning of model 1
plot(models_EN[[1]])


```



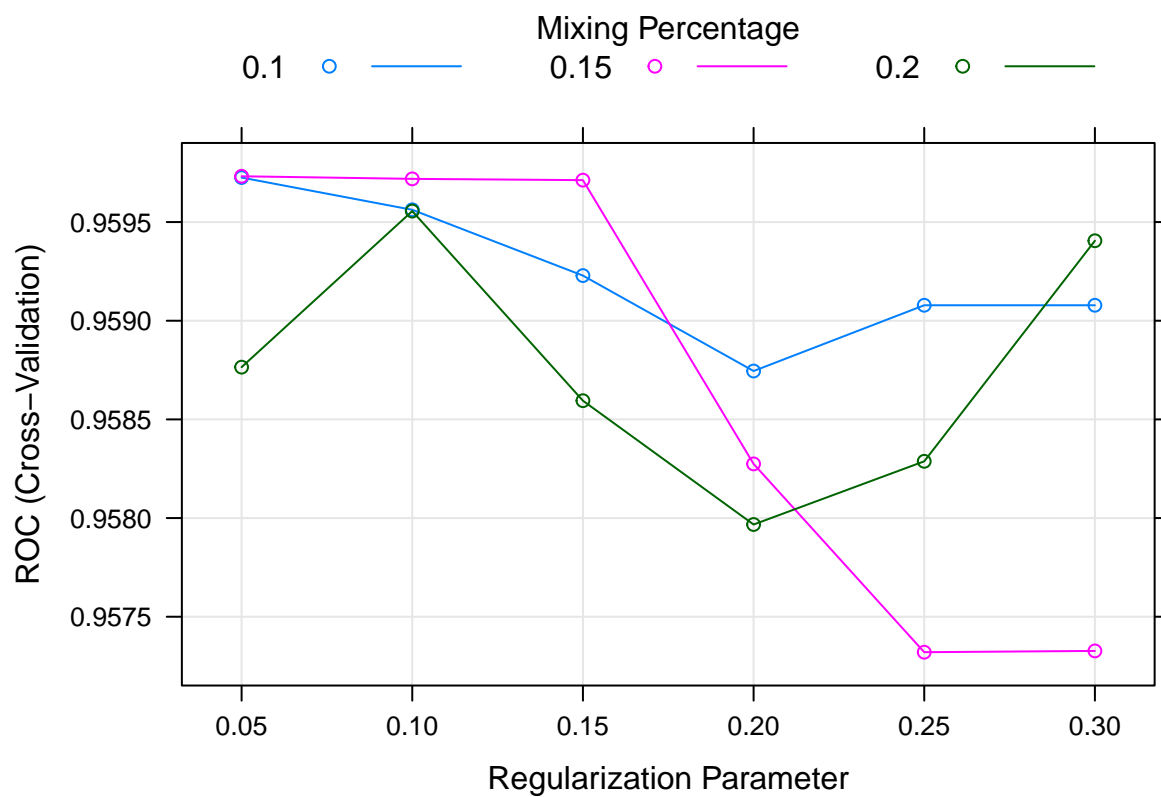
```
#Plot parameter tuning of model 2  
plot(models_EN[[2]])
```

```
#Plot parameter tuning of model 3  
plot(models_EN[[3]])
```



```
#Plot parameter tuning of model 4  
plot(models_EN[[4]])
```



Random Forest

Model comparison

Predict on the validation data set and get statistics in a table

*Select the best one at the end

Effect of class imbalance on the best model (if time allows)